

Palgo - Algorithm Animator

Andrei-Mihai Nicolae

School of Computing Science
Sir Alwyn Williams Building
University of Glasgow
G12 8QQ

Level 4 Project — March 30, 2017

Abstract

Understanding algorithms is both very common and hard for developers in general, regardless of their level of expertise. Even the fundamental ones, such as Dijkstra's algorithm for finding the shortest path between two nodes in a graph, are quite complicated to grasp. Many studies show that visualizing an algorithm and its steps make understanding it much easier. In this report, we will present an Algorithm Animator built specifically for solving this problem in a modern, responsive and efficient manner. Among others, we will also show why certain design decisions (e.g. making it a native desktop app instead of a basic jar, using material design for the user interface), the implementation choices and the evaluation results make this tool a viable option for software engineers when it comes to learning different kinds of algorithms.

Education Use Consent

I hereby give my permission for this project to be shown to other University of Glasgow students and to be distributed in an electronic format. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Name: _____ Signature: _____

Contents

1	Introduction	1
1.1	Aims	2
1.2	Motivation	2
1.3	Contributions	4
1.4	Report Content	4
2	Background	5
3	Requirements	9
3.1	Problem Analysis	10
3.2	Requirements Gathering	10
3.3	Functional Requirements	11
3.4	Non-Functional Requirements	11
4	Agile Software Development	12
4.1	Extreme Programming	12
4.2	Planning	14
4.3	Continuous Integration	15
4.4	Issues & Bug Tracking	16
5	Design	18
5.1	Architecture	18
5.2	Application Type	19
5.3	Prototyping	20

5.4	Electron	21
5.5	Drawing	22
5.6	Animation	23
5.7	Material Design	24
5.8	Compromises	25
6	Implementation	26
6.1	Algorithm Selection	26
6.2	Project Structure	26
6.3	JavaScript and Multi-Threading	27
6.4	Extra Features	29
6.5	Lessons Learned	31
7	Testing	32
7.1	Unit Testing	32
7.2	Integration Testing	32
7.3	Prototype Evaluation	32
7.4	Results	32
8	Conclusions	33
8.1	Open Source	33
8.2	Project Roadmap	33
8.3	Final Thoughts	33
8.4	Acknowledgements	33

Chapter 1

Introduction

Firstly, the starting point should be defining what an algorithm is.

NOUN

A process or set of rules to be followed in calculations or other problem-solving operations, especially by a computer.
'a basic algorithm for division'

Figure 1.1: Algorithm definition taken from the Oxford English Dictionary.

As we can see above, it is a "process or set of rules" followed by a computer when trying to solve a problem. As a machine is, at its core, composed of 1s and 0s, a human being needs to visualise what is going under the hood in order to comprehend the steps that are undertaken. In a study conducted by Christopher D. Hundhausen et al. [9], it was proven that even though algorithm visualisation is not the "perfect" solution, it is definitely effective in teaching students and experienced developers.

There are many tools out there that already provide this functionality, but one thing was a key factor in deciding how to proceed with planning the whole software process: there are barely any standalone applications. Even though one can find many websites which let people visualise algorithms' steps (e.g. a great example is VisuAlgo [17]), this is not such a good solution as the user cannot use the application at his/her commodity, an Internet connection being required.

Therefore, the plan was to create an app that can be run offline and would feel natural to the user regardless of the operating system of choice. There are many Java implementations (i.e. resulting in jar executables), but one of the main drawbacks is the lack of proper GUI tools that can build native-feeling applications.

As such, the decision was made to use the Electron framework originally built by the team behind the Atom text editor. This framework uses only web tools to generate executables for all 3 main OS families (i.e. macOS, Windows and Linux) that provide the user with a native, modern and responsive feel.

The planning was made using the best agile practices [5], eventually deciding on following a variation of XP programming. Designing and coming up with an architectural plan was a major milestone to be reached as it took a considerable amount of time to be put in place. However, as it will be discussed in future chapters, the implementation of the animation engine was a crucial and time consuming challenge that required the highest amount of effort.

The report also presents how the app was tested in many various ways, as well as how it was evaluated using potential end users. In the end, we shall discuss about the roadmap of the project, why open source is and will

be vital for the development of the animator as well as some final thoughts and lessons learned throughout the process.

1.1 Aims

The goals of the project were set and subsequently refined throughout many project meetings, as well as meeting with some fellow classmates to get feedback along the way.

The main aims of the Algorithm Animator, however, have always been:

- Create an efficient and easy-to-use animator.
- Make the animator a cross-platform application that would run natively on the main operating system families: macOS, Windows and Linux.
- Provide a user-friendly interface that would make the user want to enjoy the product.
- Create at least 5-6 fully functional algorithms.
- Make the application scalable and easy to maintain.
- Test all functionality and ensure quality above quantity.
- Evaluate the product throughout the software process to meet acceptance criteria.
- Build a roadmap that would allow the animator grow even after the level 4 project has finished.

1.2 Motivation

Even if we have previously mentioned that visualising algorithms is one of the best ways of understanding how they perform, there are also other key motivational aspects behind the need of building an animator.

One such motivation is the increasing demand of software engineers throughout the industry as well as the rise in level of expertise. Developers, and students in particular, will benefit tremendously from a good grasp on how to implement correct and optimal algorithms when applying for a new job. Thus, an algorithm animator would be a good component in one's tool belt.

Another major aspect was the opportunity to create a modern tool that would make users enjoy working with it. At the moment, one of the most popular animator toolkits is the one presented in John Morris' paper [20]. Below a screenshot was taken from the app.

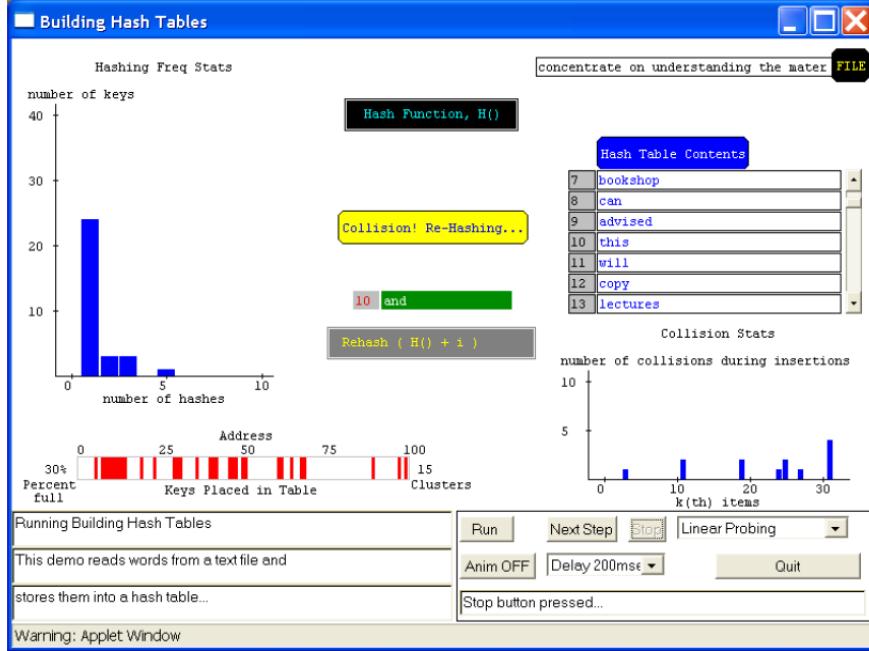


Figure 1.2: John Morris' Animator Toolkit.

It was written in Java, which comes with certain advantages: multi-threading support, enhanced familiarity due to university coursework and projects etc. However, the lack of proper graphical user interface components makes this a not so viable option for regular users when adopting it.

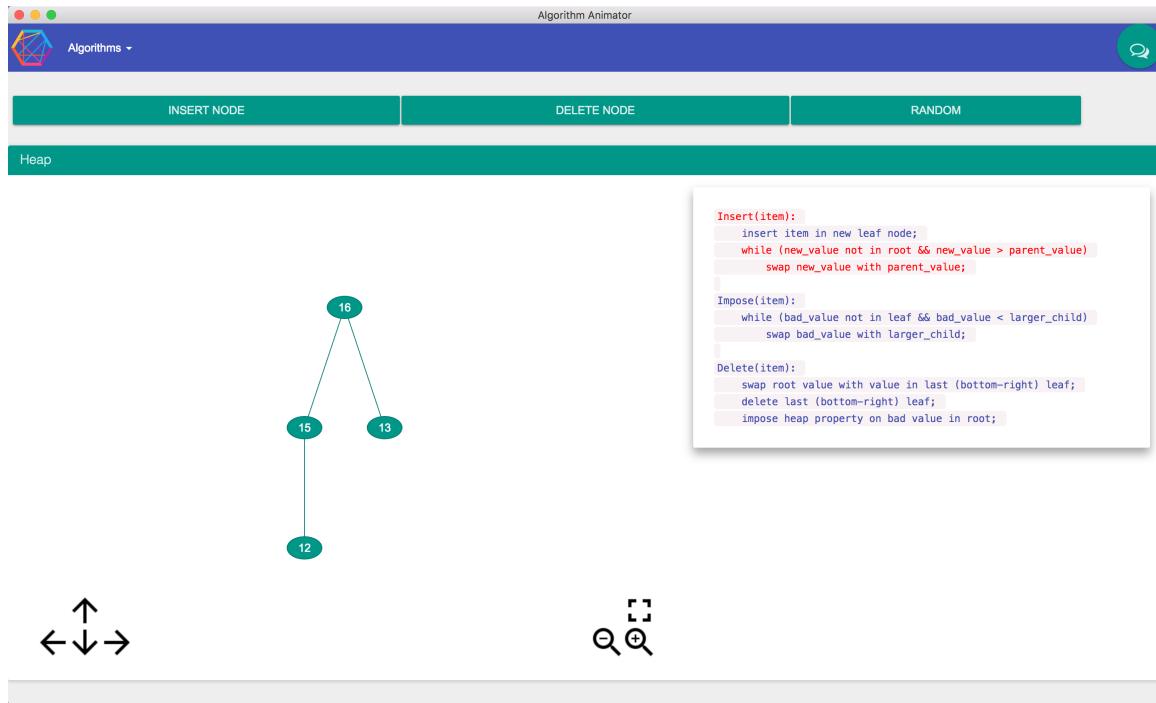


Figure 1.3: The animator presented in this report.

We can see a clear difference between the two and why one might choose the latter option due to familiarity (native feel) and easy-to-use user interface.

1.3 Contributions

This report serves to deliver certain main contributions:

- Shows the whole software development lifecycle of an algorithm animator.
- Presents various technologies used to build native and modern desktop apps [12] [32]
- Provides a brief overview of current animators available and background knowledge on their benefits and efficiency in teaching.

1.4 Report Content

The rest of the report will analyse the background of animators and why they were proven useful, as well as cover all the steps in gathering requirements, designing, implementing, testing and evaluating the tool.

- Chapter 2 covers work related to the purpose of algorithm animators and why they are useful
- Chapter 3 goes into how the problem was analysed and what requirements were gathered through project meetings and discussions with Algorithmics students.
- Chapter 4 shows the steps undertaken to follow the best agile methodology principles.
- Chapter 5 explains the design decisions behind the tool and illustrates various lessons learned and problems faced along the way.
- Chapter 6 goes into the implementation details of the animator.
- Chapter 7 show how extensive unit, integration and other types of testing (e.g. smoke, end-to-end) were undergone and why they were essential to the development of the application.
- Chapter 8 details the overall results of the project.

Chapter 2

Background

Firstly, we need to define an algorithm. We will use the above-mentioned definition taken from the Oxford dictionary [10]: "A process or set of rules to be followed in calculations or other problem-solving operations, especially by a computer". In our case, we will only relate to set of rules followed solely by computers. Thus, put simply, an algorithm is a set of steps defined by someone in order to solve some problem. As an algorithm can be written by anyone, there is no known number or catalog of them. However, the most important algorithms (e.g. Huffman tree encoding/decoding) are taught in any Computer Science course at universities around the globe.

Having defined what an algorithm is, let's take an actual example - Dijkstra's algorithm for finding the shortest path between two nodes [11], first published almost 60 years ago (the pseudocode was formally checked by Dr. Norman):

```
1 // S is set of vertices for which shortest path from u is known
2 // d(w) represents length of a shortest path from u to w
3 // passing only through vertices of S
4 S = {u}; // initialise S
5 for (each vertex w) d(w) = wt(u,w); // initialise distances
6 while (S != V) { // still vertices to add to S find v not in S with d(v) minimum;
7     add v to S;
8     for (each w not in S) // perform relaxation
9         d(w) = min{ d(w) , d(v)+wt(v,w) };
10 }
```

Listing 2.1: Pseudocode for Dijkstra's shortest path algorithm

As one might observe, this is not trivial to grasp. Therefore, many solutions have been created in order to help students and engineers in general understand them better. One such solution is a rather creative one, developed by students at Spatienta University, which tries to teach different sorting algorithms by dancing.




Figure 2.1: Sorting algorithms learning through Hungarian folk dancing.

This was featured on many websites and gained enormous traction among computer science enthusiasts. However, a more traditional approach that is adopted formally in universities as well is the use of animators.

Firstly, let's define what an animator is. Again, taking the definition from the Oxford dictionary [10], it is "a person who makes animated films". In our case, the person is the computer and the film is the sequence of steps required to produce an output after feeding input to an algorithm.

Usually, most algorithm animators are divided into 2 sections: one side will contain the graphics in order to represent the data structures used (e.g. in graph traversal algorithms the nodes and edges comprising the graph will be drawn) while the other side will have the code lines, the actual rules the computer needs to follow in order to produce the desired output. In general the current code line that is executed at a certain moment is highlighted, while the data structure visualized (e.g. some box that represents an array's element) changes and the user can see the output visually and immediately, making a mental connection between the code line and the change it produces.

But why an animator precisely? There are many studies that show visualization in any area can enhance learning capabilities drastically. Brenda Parker and Ian Mitchell's work [24] shows clearly how seeing the steps can greatly improve a student's understanding of what is happening "behind the scenes". On the other hand, even if animations were proven to be generally useful and beneficial for the learner, the need for seeing algorithms animated declines over time, as the study shows.

Another study by Susan Palmiter and Jay Elkerton [23] tried to find out which method was the most effective in learning new computer-based tasks. Even though animations for complex mechanisms were found not to be as effective as for some easier one (e.g. parallel programming vs. graph traversal), they are very helpful for the algorithms generally taught at university courses or used most often.

Having discussed methods of learning new algorithms better and why animators are useful, we need to take a look at previous solutions developed by other people. Apart from John Morris' tool shown above, there are many other applications available, including Galant [29] and the visualizer built by D. Galles [13].



Figure 2.2: Galant Algorithm Animator.

Here we can see one of the 2 main approaches of building an animator, which is a standalone JAR application. Thus, it is written completely in Java and, as the language is pre-installed on most electronic devices, it can run on the major operating systems families.

There are obvious advantages to this approach:

- end product can run on a variety of machines with only one codebase
- large number of Java libraries facilitate an easy development workflow
- many tutorials and resources available
- java is taught at most universities in introductory courses, making it easier for newcomers to create their own application rapidly

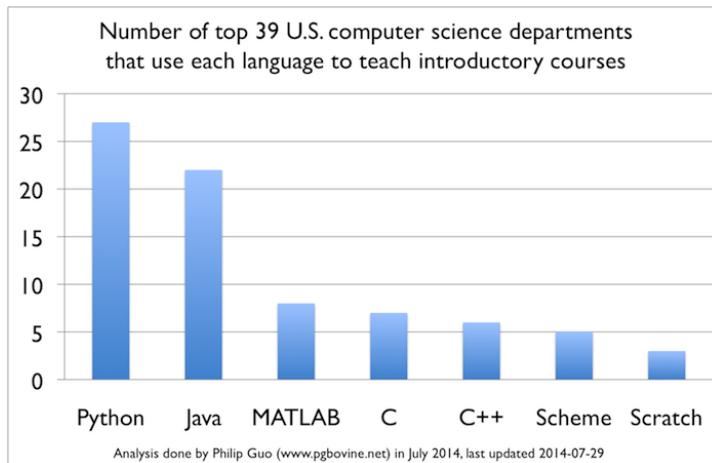


Figure 2.3: Languages used in introductory courses at top US universities.

However, there are also drawbacks:

- the most popular Java libraries for producing standalone graphical JAR applications output old-fashioned user interface
- the code for drawing on the canvas and animating in Java (using Swing or JavaFX) can become extremely verbose and unreadable

On the other hand, we also have the web approach to building animators, one of them being D. Galles' application [13].




Figure 2.4: Algorithm animator built by D. Galles at University of San Francisco.

Some advantages of this solution are the following:

- the developer can use a wide variety of JavaScript tools (e.g. drawing directly on HTML5 canvas, VisJs, SigmaJs) to draw and animate objects
- the app can have a very modern and responsive feel which can make users enjoy the product
- virtually all users have at least a browser installed on their machine, thus making the animator available to everyone

However, the main drawback is that it is still a website, thus it need to be hosted (which implies additional costs) and it cannot be reached by anyone at all times due to the fact that not everyone has constant Internet access.

We have covered all the necessary background and related work before diving into how Palgo was designed and implemented. We believe that our solution is a combination of the two main approaches listed above, merging the benefits of both into a single tool.

Chapter 3

Requirements

As for every software built with best practices, requirements need to be gathered, defined, analyzed and then prioritized. As we tried to follow an Agile structure, the requirements have changed every week through student-supervisor weekly as well as end user meet-ups.

We tried to keep the requirements SMART [19], which implies they are specific, measurable, attainable, realizable and time bounded. Thus, making them follow this well established format made the software development lifecycle much faster and cleaner.

Below is a screenshot from a Trello [30] board used during one of the iterations (a separate board was used for tracking non-functional requirements as well).

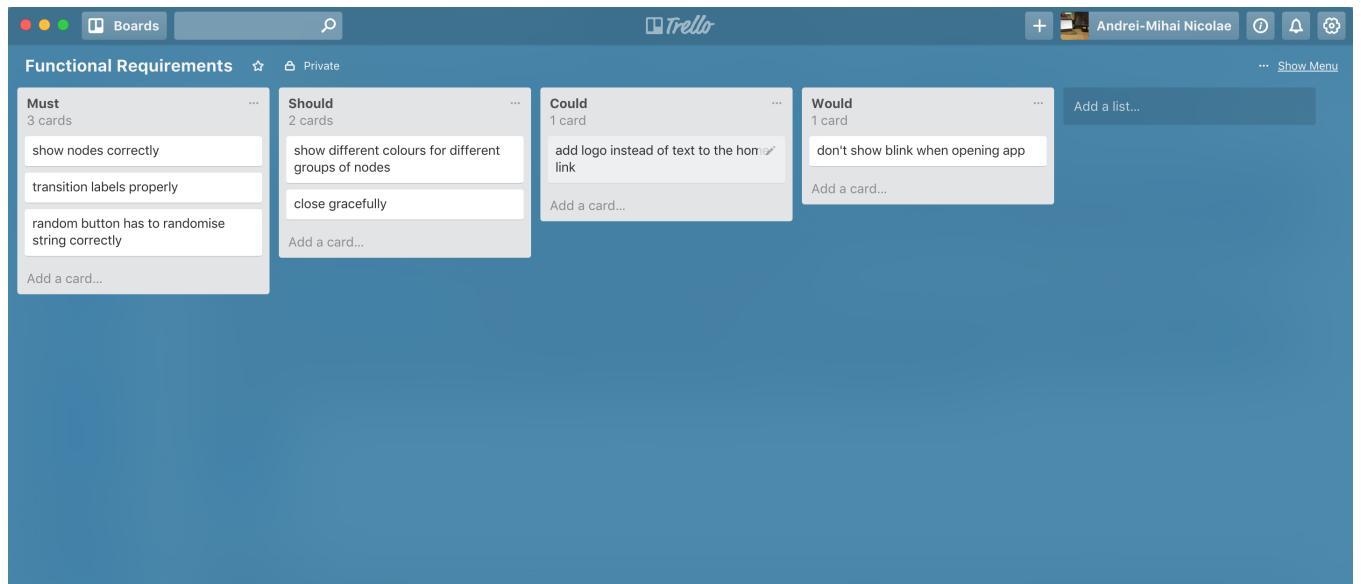


Figure 3.1: Trello board for functional requirements during one of the iterations.

The requirements followed a pipeline through which the requirement together with its rationale would get broken down into pieces continuously until user stories would be produced. One example of such a requirement would be: "Show me a correct and easy-to-understand animation for the insert function in a heap data structure.". This would then be broken down in multiple user stories, such as "Color current node that is inserted and add its label" and "Highlight code line corresponding with changing the label".

We also tried to follow the MoSCoW [7] guideline for prioritizing the requirements. The technique has

proven to be very effective and it is adopted widely in industry, even though the paper is more than 20 years old. This facilitated a smooth workflow and concise goals to be targeted.

3.1 Problem Analysis

We have analyzed the problem systematically, devising it into multiple main questions to be answered:

- who are our target users?
- what type of format will the application have (i.e. web/desktop)?
- what algorithms should be animated?
- how would the main layout be structured?

During the analysis, we needed to take into account various factors, including how much experience with algorithms and computer science in general will our users have or what algorithms would need visualization the most. As previous animators built under the supervision of Dr. Norman are used for educational purposes in the level 3 Algorithmics I course, we believe that having mainly graph and tree algorithms would benefit the project the most. However, we tried to keep the structure simple and the codebase easy-to-maintain such that other families of algorithms will be straightforward to implement. Because we want to project to be continued in the open source space, users other than students can use the animator, thus we tried to keep the requirements more general.

Also, there was much discussion around whether to use a web or a desktop approach. For reasons explained above, we resumed to implementing a desktop application using the technologies mentioned in the subsequent chapters.

We have gathered requirements and split them into functional and non-functional ones, dividing them further into smaller tasks to be implemented. At the end, we were left with user stories that could be translated directly into work tasks.

3.2 Requirements Gathering

As mentioned above, the requirements gathering was a continuous process that lasted from first week of 4th year up until the end of the project's development lifecycle. Not only that we had the weekly regular student-supervisor meetings where we could add or refine requirements, but there were also meetings with possible end users.

Every two or three weeks there were meetings scheduled with a couple of classmates from University of Glasgow's Computer Science course that would give continuous feedback after testing the newest version of the application. We met in one of the laboratories in Boyd Orr building where they were provided with an executable that could be run on their personal laptops. They were given around 15 minutes to try the features and, afterwards, they provided invaluable feedback without which the animator would not perform as efficiently as it does today.

3.3 Functional Requirements

Functional requirements target specific functionalities that the system should perform. They started with a couple of algorithms that needed to be animated. Then, functionalities such as current code line highlighting and specific button behaviors began to stack up. Eventually, we can categorize them as follows:

- animate tree/graph depending on what the user chose to do (e.g. insert/delete node)
- highlight current code line for every algorithm
- define specific behavior for every button displayed on the UI
- each window interaction should perform as expected depending on the operating system the user is using the application on

Many smaller ones appeared along the software process, but the above mentioned are the main categories.

3.4 Non-Functional Requirements

A non-functional requirement is a type of requirements that judges the system as a whole and how it operates rather than checking for specific criteria. The main non-functional requirements were rather set from the beginning. We decided to have an algorithm animator that would:

- be easy to use and let the users familiarize rapidly
- make users enjoy playing with the product
- be installed/fetched with ease
- allow end users to use it on any platform
- be efficient and not consume excessive resources
- be maintainable and scalable (i.e. let future possible contributors add algorithms easily)

Chapter 4

Agile Software Development

Good software practices are key to successful projects. We have followed Agile methodologies, especially what was taught last year in the Professional Software Development class, as well as through Ian Sommerville's "Software Engineering" book [27].

We eventually decided to follow an Extreme Programming path and we implemented the best advice taken from industry experts. Some of the key principles we implemented throughout the whole process are:

- good planning strategy that would allow us to follow a schedule; our cost estimation efficiency improved with every iteration
- frequent commits and releases that would generate continuous feedback from our potential end users
- customer involvement (i.e. student meetings every couple of weeks as well as weekly student-supervisor meetings)
- heavy refactoring as opportunity arose
- emphasis on simplicity in design and implementation
- tracking bugs and issues in general allowed us to be always informed and include necessary items for specific milestones so that the quality would always have high standards

All these ensured a good workflow with results that could be remarked immediately. We believe that a major part of the success of the project was embracing the above Agile methodologies.

4.1 Extreme Programming

We have followed the approach first coined by Kent Beck called Extreme Programming [3]. It is an Agile technique through which good practices are most times pushed to "extreme" levels. XP has proven to be very effective as many companies in the industry use it in their development lifecycle. Some principles it promotes which we have followed are:

- frequent releases based on requirements gathered incrementally in meetings (e.g. supervisor-student meetings)

- constant refactoring to promote simplicity
- customer involvement throughout the whole process (i.e. getting feedback every 2-3 weeks from students that might use the application)
- people over processes, which means collective ownership of the code as well as avoiding exhausting long working hours

Burndown charts are another great tool used by people using XP in their schedule to track down progress. You start off by defining the number of user stories you want to complete in an iteration and, then, divide that number by the number of days the sprint has (in our case it was 7 days). Therefore each day you should implement a certain number of user stories, constantly, such that at the end of the iteration you are left with none. That is the most productive outcome. Also, because this was a student project, we will count weekend days as working days.

During the entire process we tried to monitor the progress and see how we are performing. Below you can see a burndown chart from one of the November weekly iterations.

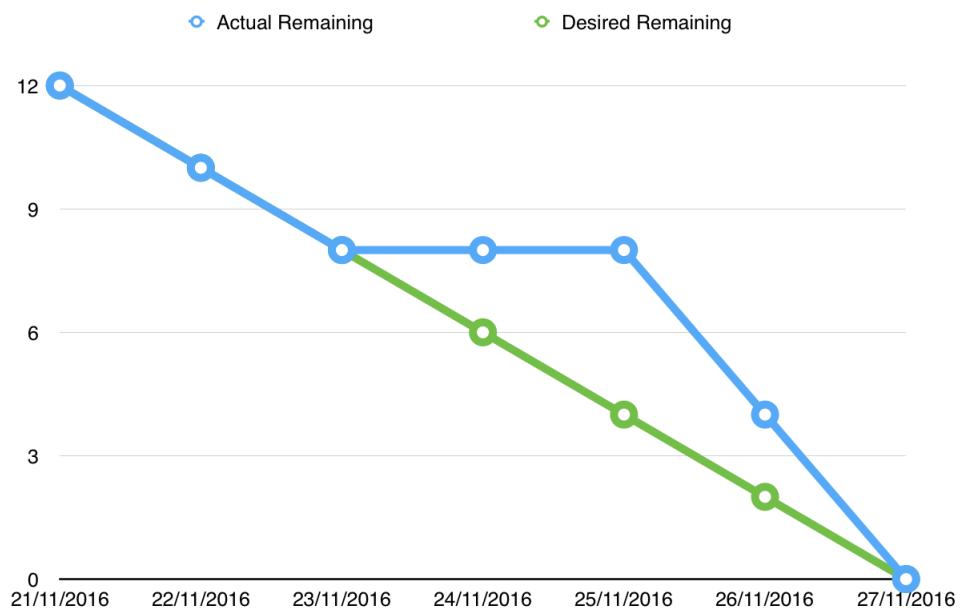


Figure 4.1: November iteration burndown chart.

As one can notice, there is quite a discrepancy between the desired number of user stories per day and what we actually completed. This was because of lack of consistency, shown in the graph between days 23 and 26, where there were no user stories finished. However, below is a burndown chart corresponding to one of the last iterations of the project.

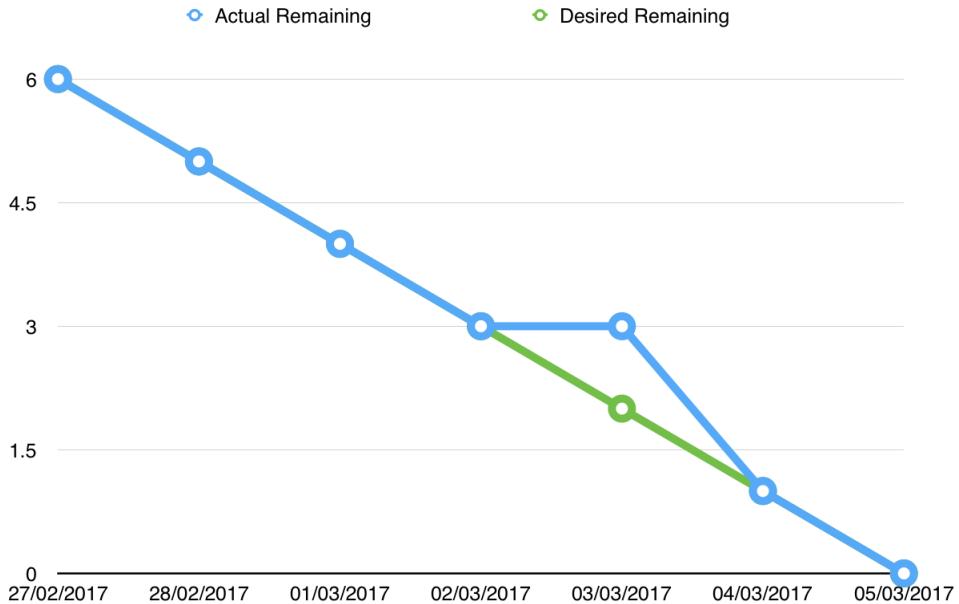


Figure 4.2: February-March iteration burndown chart.

There are noticeable difference between the 2 iterations: firstly, we started with fewer user stories to be completed in the latter sprint which was much more feasible to finish and it didn't make us work extra hours (like in the first burndown chart example), thus violating one of the extreme programming principles. Moreover, there was only a slight discrepancy from the schedule on March 3rd where there were no user stories completed. However, this is relatively minor and we believe that for a two people project, this was an amazing achievement that could only be reached gradually.

Summing up, following all the good practices of XP has brought us great benefits and a working schedule that allowed us to implement all the requirements gathered in user and supervisor meetings.

4.2 Planning

Planning is a vital part of every Agile project - it involves setting deadlines for tasks needed to be implemented by developers. However, as some tasks are not trivial to estimate, multiple strategies have been invented so that team members can effectively set number of days (or other units of work) for every user story.

In our case we adapted the most effective Agile method, which is Planning Poker [15]. In this practice, all developers are given decks of cards and when a user story is brought up, every developer shows a card which has a value. That value will represent the number of units of work needed to complete the story. Then, after everyone shows their estimation, discussions emerge among the developers until a consensus is reached. In our situation, we did not use a deck of cards during the student-supervisor meetings, but we both estimated the number of days needed to implement a feature and, if our opinions diverged, we discussed until we had the same duration in mind.

Below it can be observed how we kept track of all the user stories using Trello cards. We added the estimation in the description, as well as a description using the popular format "As a.../I want to.../So that..." [8] to always know the rationale behind every implemented feature.

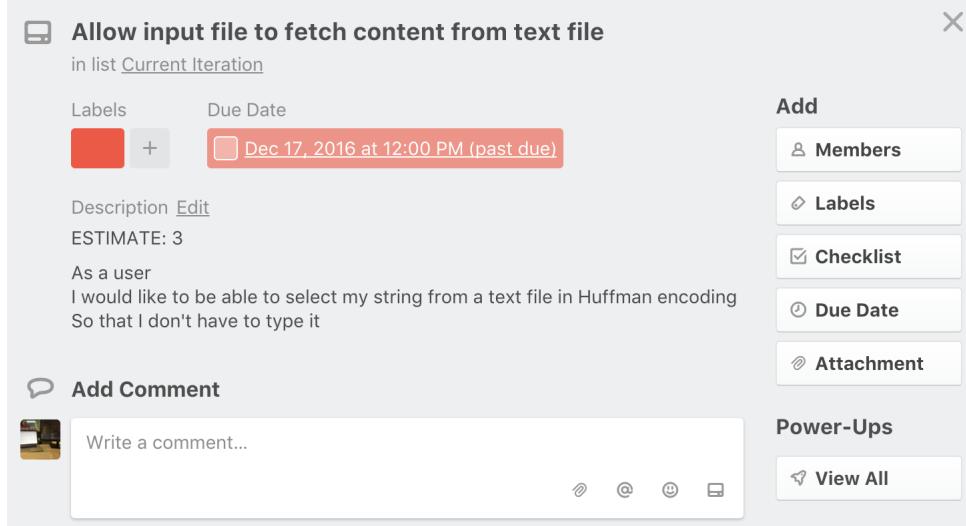


Figure 4.3: Trello card for estimating a user story.

This method was very effective and as we gained more experience after several iterations, we would come to a conclusion much faster and the user story cost estimations became more accurate.

4.3 Continuous Integration

We have used another great tool heavily promoted by Agile methodologies which is Continuous Integration. CI means merging all working copies of everyone who is developing on the project to a shared mainline multiple times a day [4]. This is an amazing tool for checking that the builds do not break regardless of the changes made and that quality is of primary importance.

In our case, we have used CircleCI [6] as the tool for CI. It was configured such that after every commit, the code would be merged into the main branch, tested to see if everything is working as expected and, only after, the code would be pushed.

 removed unnecessary edge popup for heap nclandrei committed 8 days ago ✓	View	Copy	aaefac6	Compare
 changed layout of heap algorithm page nclandrei committed 8 days ago ✓	View	Copy	1db77c1	Compare
 refactoring in prim jarnik versions nclandrei committed 8 days ago ✓	View	Copy	delf1d1	Compare
 solved bug that did not remove items from ntv table nclandrei committed 8 days ago ✓	View	Copy	a2aaa25	Compare

Figure 4.4: CircleCI status showing red/green depending if the build was successful on GitHub.

As it can be seen above, there is a small green check (or a red x if the build did not pass) on the GitHub repository of the project telling the developers if anything went wrong. This was a great asset to the development lifecycle, informing us whether we needed to add some bugs to our tracking system to be solved later or everything performed exactly as it should have.

Not only that we applied continuous integration to our workflow, but we also used it properly: as mentioned above, the definition of CI means the developers should push their local changes to the main repository multiple

times a day. Every day of development meant multiple commits and multiple pushes such that we made sure the builds were all passing and that our code was not broken.



Figure 4.5: Commit frequency in the last 10 sprints.

This figure shows the commit frequency per week for the last 10 iterations. As it can be observed, apart from first week when the university semester just started, the number of commits was very high every sprint, thus following the practice of committing and testing the quality of our product often.

4.4 Issues & Bug Tracking

Issues and bugs occur all the time in every piece of software. Because the userbase was rather limited, I updated the database manually from what we (i.e. Dr. Norman and myself) found not working as expected and what the students marked as functionality not working properly or wanted features. This was a very useful practice to follow because it also promoted continuous feedback on changes from potential end users, thus implementing the Agile manifesto properly.

Talking about actual implementation, we have kept a custom database on GitHub in order to be always up-to-date with what is going on and spots where we lack quality and efficiency. We have made use of the custom labels GitHub already provides, among which the most used ones have been:

- bug - here we marked flaws in the system (or the students who have been providing constant feedback)
- enhancement - we added here the features most wanted by our users or by ourselves; this was refined from the requirements gatherings previously mentioned that ended up as multiple user stories eventually
- question - these were marked for research purposes - one example might be: can we change the shape of nodes dynamically as the algorithm animation makes progress?

As an example, the following figure will show how issues were continuously tracked throughout the software process. This screenshot was taken during one of the iterations.

The screenshot shows the GitHub Issues page with the following details:

- Header navigation: Code, Issues (4), Pull requests (0), Projects (0), Wiki, Pulse, Graphs, Settings.
- Search bar: Filters (is:issue is:open), Labels, Milestones, New issue button.
- Issue summary: 4 Open, 1 Closed.
- Issues listed:
 - #5 Add play/pause/next/previous buttons to the algorithm animation. enhancement. opened 18 seconds ago by nclandrei.
 - #4 Get the animated nodes for the Prim Jarnik algorithm display proper colors. bug. opened 2 minutes ago by nclandrei.
 - #3 Add upload text button for the Huffman tree. enhancement. opened 3 minutes ago by nclandrei.
 - #2 Random button for Heap is not working properly. bug. opened 3 minutes ago by nclandrei.
- Filtering and sorting options: Author, Labels, Milestones, Assignee, Sort.

Figure 4.6: Example of issue and bug tracking in the GitHub provided Issues page.

As it can be seen in the image, there is also a milestone tab assigned to each issue. Because the screenshot is from an earlier sprint, there were not milestones attached, but since January every issue had one so that we would plan effectively and reach all goals by the negotiated deadline.

Chapter 5

Design

This chapter will cover the design decisions taken throughout the software process. It will delve into the architectural approach we followed, why the main technologies were chosen eventually (e.g. Electron, Vis.Js) and why they proved to be the most effective to achieving our purposes. In the end, we will also discuss compromises that needed to be taken and why.

5.1 Architecture

Because we would have a wide variety of users (students, experienced developers, university lecturers, enthusiasts) one needs to take into account a range of different perspectives. Therefore, architectural design and focus on simplicity is vital to the success of every software. We have tried to use UML diagrams in order to visualize the system as well as possible before proceeding with the implementation.

Here is our final candidate for the component diagram.




Figure 5.1: UML component diagram.

It can be noticed that our approach was rather simplistic - after the user hits the start algorithm button, the main component is the Electron window. This is, in fact, the native window generated by the framework

which contains the app. As we will see in the subsequent parts, Electron uses web tools to generate the output, thus we start rendering on our HTML page the animation. This will be composed out of 2 components: the node manipulator which will color, move, hide and change the nodes in the graph or tree, while the codeline highlighter will handle the animation of the algorithm's codelines. They will execute concomitantly and generate output immediately in the window.

Therefore, our architectural style of choice is component-based: every algorithm is a component itself and they all interact with the Electron core functionality such as windowing, copy-paste-cut, system calls etc. It emphasizes on reusability through loosely coupled independent components, thus adding simplicity to the project and making it scalable and maintainable.

5.2 Application Type

This was a very important question right from the beginning. Having the experience from past years students who implemented algorithm animators, creating a Java application that could be packaged into a JAR seemed attractive. Java is the language used the most for university coursework, thus there was comfort in choosing this approach.

Also, there are many projects available that were of interest. The one we studied closely before deciding on whether to choose Java or not as the primary language behind our animator was ANIMAL [26], one of the most popular tools on the web. Even though the paper published by Rößling et al. is almost 20 years old, the animator is still actively under development, the latest release dating back to 5th of September 2016.

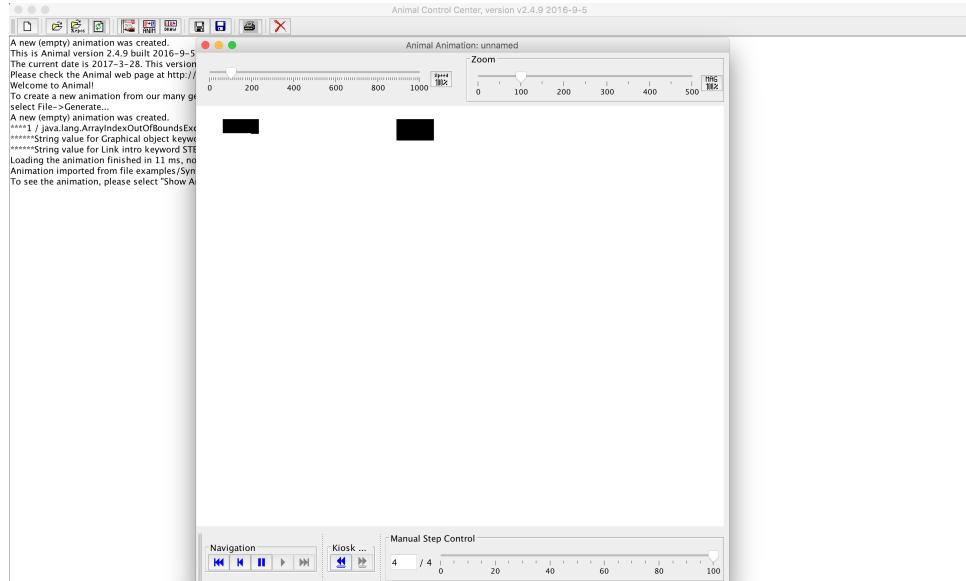


Figure 5.2: Screenshot taken while using the Animal algorithm animator.

This screenshot shows how many options are available to the user and how powerful the application is, providing features such as the possibility to create your own algorithm, animate it the way you want to, set any color for blocks and even draw manually.

On the other hand, the other option that we considered was making an animator based on web tools. The first and obvious approach was to create a website using the very large number of libraries for drawing on the canvas and manipulating the DOM in general. Firstly, we began by investigating how popular JavaScript was among developers and if this was feasible after all. Here is a brand new survey from StackOverflow, one of the most comprehensive and useful resource repository for developers around the globe [28].

Programming Languages




Figure 5.3: Most popular 18 programming languages in 2017.

As it can be seen, JavaScript is with more than 50% popular than its competitor Java. Also, when talking about concrete JavaScript libraries, there is a considerable amount published as open source. Other advantages for this approach were:

- great reach to users as almost everyone has at least a browser installed on their machine
- modern, state-of-the-art UI libraries that would make the experience pleasant for users

Weighing both options, we eventually decided to proceed with the web approach because our main goal was to make users, especially students who want to learn new algorithms, like using our animator. Web tools, as discussed previously, provided us with a variety of User Interface components that made our goal possible.

5.3 Prototyping

TODO: include here low fidelity and high fidelity prototypes.

5.4 Electron

Thus, we ended up with using a web approach. There are many options available, ranging from university-specific algorithm animators to more general ones such as the previously mentioned VisuAlgo [17].

However, one very important aspect that we were concerned about was the core underlying principle of the internet, which is that it requires a permanent internet connection. What if our users would not have access to internet a couple of days before an algorithmics exam? What if our potential user would want to play around with algorithms at his/her own pace without needing to have a wired or wireless connection?

Also, another thing we had to think about was UI and UX engineering. Having maybe too many resources available, we could have ended up with choosing one that would not fit a part of our users. How can we make sure that the framework we would decide on will make all users feel familiar?

Therefore, the best approach would be to have an animator that is implemented using web tools but it can function as a standalone desktop application. Thus, by being native, it would make all users feel familiar in any their operating system of choice. At this point, we had to research different alternatives to create cross-platform native apps with web tools. The one that was most popular and had the largest library of resources was Electron, which we chose in the end.

Electron is a framework built by engineers at GitHub and then released as an open source project on their platform. It is in fact a Chromium backend that serves webpages inside windows which act as native regardless of the platform of choice. Therefore, the developer can use any web tools available to develop his/her applications benefiting at the same time from features such as system calls.

Because it is using Chromium under the hood, we are automatically taking advantage of using the V8 engine [14] which powers the Google Chrome browser. Efficiency was primordial to us and such we wanted to check if such a candidate would produce expected output in a reasonable time, so we checked benchmarks available for various JavaScript engines.




Figure 5.4: Benchmarks for various JavaScript libraries as of 28/03/2017.

As we can see above, the V8 ranks extremely good, being at the lowest points of execution times starting almost from its inception.

The other very important aspect was to see if Electron had enough resources, tutorials or demo applications such that the development workflow would be smooth. We ran a query on Google Trends among other platforms to check how popular it was and we were given the following results:




Figure 5.5: Electron trending on Google in the past 5 years.

We can clearly observe that apart from the big spike when the framework became the underlying engine behind many important desktop apps such as Slack somewhere at the end of 2015, Electron has had a constant ascension throughout the years. Also, researching the internet for resources, we could find many tutorials and demo apps where we could draw inspiration from.

Summing up, Electron was a natural choice after careful thinking about advantages and caveats and it proved to be a very productive framework that allowed an easy and simple workflow for the rest of the software development lifecycle.

5.5 Drawing

The component that had the same importance as Electron was the engine behind manipulating the nodes and edges of tree/graph algorithms that we wanted to implement. Because of the large variety of JavaScript libraries available for playing with the canvas, we had a couple of choices to select from. This was a very lengthy process because many provided some features while others gave us other unique ones.

We shall list them below, together with their pros and cons:

	Nodes Manipulation	Edges Manipulation	Tree API	Graph API	Actively Maintained	Easy Pick-Up
Treant	✓	✓	✓	x	x	✓
Vis JS	✓	✓	✓	✓	✓	✓
Sigma Js	✓	✓	✓	✓	✓	x
HTML5 Canvas	x	x	x	x	✓	x
D3	✓	✓	x	x	✓	x

Table 5.1: Table showing features provided by the main drawing libraries for trees/graphs.

These are the main candidates we have chosen: VisJS, Treant, D3, HTML5 Canvas and SigmaJS.

Let's start with Treant. This was a very compelling option because it had a very robust and easy to use Tree API: nodes and edges could be easily manipulated, as well as the layout was very flexible to suit the developer needs. However, as it can be seen in the table above, we did not have an out-of-the-box API for graphs. Therefore, this option became rather unfeasible.

HTML5 Canvas is a component defined by W3 Consortium themselves and thus it's actively under development and heavily standardized. However, apart from continuous maintenance and new regular features, it is hard to pick up and it does not provide APIs specifically designed for graphs or trees which meant we needed to put too much effort in learning a new platform rather than focusing on actual implementation of the animations.

D3 and SigmaJS are both very mature platforms that have a very large number of features for animating and drawing on the canvas. They were both viable candidates up until the end. However, the reason why we eventually chose VisJS is because it was the easiest to pick up and had out-of-the-box APIs for both graphs and trees.

Thus, choosing VisJS, we saw a rapid development workflow and an emphasis on simplicity due to the fact that the tools we were using were not standing in our way, but rather helped us to quickly implement the user stories.

5.6 Animation

After drawing, the next major issue we had to discuss was the actual animation part. We have researched many tools available and we ended up with basic set timeouts. One snippet for an animation looks something like this:

```

1 setTimeout(
2   function() {
3     unHighlightAllCodeLines();
4     for (var k = 0; k < nodesArrayLength; k++) {
5       unHighlightTableCell(nodes[k].label);
6     }
7     resetWholeNetwork(network, container, options);
8   },
9   2000 + 3000 * nodesArrayLength + 13000 * nodesArrayLength - 1
10 );

```

Listing 5.1: JavaScript snippet for creating an animation

Even though the implementation details will be covered in the following chapter, we shall discuss here the importance and the reason why we decided to use this approach instead of others.

Some libraries were already built to facilitate developers an easy path to animation, so why use plain JavaScript for accomplishing this? We have reviewed and discussed advantages and disadvantages for the following JS libraries:

- Velocity JS
- DOM JS Animation
- JQuery Animation

Starting with Velocity JS as it is an actively developed library with more than 39 contributors on Github and almost 700 commits [31], what it primarily offers is transforms, loops, easings or color animation. This sounds promising, but the main problem in our case is that we are working with Vis Nodes in a tree or a graph, thus our components are JavaScript objects rendered in the canvas. Even more than that, we have circular dependencies, as the root of a tree can have children which in turn are Nodes themselves, so mixing the structure Vis JS imposes with the functionalities of Velocity JS was rather impossible.

Also, another issue with using any of the 3 options (W3 DOM Animation, Jquery animate or Velocity) was that they are manipulating elements which are part of the DOM (Document Object Model), while in our project the elements are part of the canvas. Therefore only the canvas can be animated as a whole without being able to manipulate the elements inside.

Having discussed all these caveats, we were left with no other choice than devise algorithms to animate our nodes and codelines using plain setTimeout calls throughout the project. Even though it was a compromise and it slowed drastically the workflow in the beginning, once we passed the steep learning curve, we were able to animate all the elements in short periods of time.

5.7 Material Design

We had a couple of main possibilities to choose from when designing the application's User Interface:

- Material Design by Google
- Bootstrap by Twitter
- Pure CSS

There were others as well, such as Base [2] or Simple Grid [16], but these were not specifically suitable for our purposes. Therefore, we shall analyze different benefits given by the 3 above mentioned platforms and why we eventually ended up with Material Design by Google.

Below we shall present one of the most used and noticeable components to a user when he/she interacts with an application, which is the navigation bar.



Figure 5.6: Bootstrap default navigation bar.



Figure 5.7: Material Design navigation bar.



Figure 5.8: Pure CSS default navigation bar.

We can clearly notice a more colorful, natural and playful approach in the Material Design navigation bar. We have presented all 3 options to the level 4 students and they were all in favor of the Material Design UI for all components (more details on the results of these evaluations in Chapter 7, section Prototype Evaluation).

Even though the Material Design components were invented by Google (which are also available on GitHub), we have opted for the ones open sourced by Federico Zivolo [33] because they are much more popular and there are more resources with which one can get accustomed quicker.

We have decided to use this approach in spite of the others because of several reasons:

- Material Design has gained much more traction than the other frameworks since its inception
- even though Pure CSS does not require any JavaScript dependencies (hence the name), it is very raw and it requires a lot of tweaking before going into production
- Bootstrap is the foundation of virtually every framework available but it is exactly that, a foundation, thus it does not give the developer advanced features out-of-the-box

5.8 Compromises

One of the compromises mentioned previously as well was that we had to use `setTimeouts` to animate our objects instead of mature, easy-to-grasp libraries. What the `setTimeout` function does is it sets the execution of the callback function provided at a specific time in the future. Because we had only a canvas and we couldn't manipulate nodes as regular DOM objects, all the node coloring or label changing was done using a chain of `setTimeouts` that, at the beginning of the algorithm (i.e. when the user presses the "Start Algorithm" button), they would have their exact time in the future when they got executed calculated. One example might be: color this node at millisecond 2600, change its left children's label to 2 at millisecond 3500 and finally highlight the codeline "set label to 2" at the same time, 3500. This was extremely hard to understand and master and the first iterations were quite slow. However, as time passed and I gained more experience with this technique, development became much easier.

Another compromise that we needed to take was the lack of play/pause and previous/next buttons for the animator. Here is a screenshot taken while using VisuAlgo:




Figure 5.9: VisuAlgo play/pause/previous/next buttons for the animator.

This is a functionality desired by most of our users as well as Dr. Norman. However, due to the previous compromise we have stated (i.e. using `setTimeouts` for animating the algorithms), it was virtually impossible to mimic these buttons.

Chapter 6

Implementation

In the following sections we will discuss among others about how the project was actually implemented, how we overcame the JavaScript shortcoming of not having multi-threading support, as well as how other various functionalities were put into place. We will conclude with talking about biggest issues faced and what lessons were learned.

6.1 Algorithm Selection

The algorithm selection was solely done during the supervisor meetings. We have come up with the following list:

- Huffman tree construction
 - Depth-First Search in a graph
 - Breadth-First Search in a graph
 - Heap insertion and deletion
 - Dijkstra's Algorithm for Shortest Path
-
- **6.2 Project Structure**

We have split our application into 3 main components, as in traditional website development: stylesheets, JavaScripts and HTML5 pages. They were all coordinated by npm which is the node package manager [22]. Essentially, npm is the equivalent of Maven or Gradle for Java projects and that is both a dependency manager and a build automation tool. Therefore, in npm we defined the name of the project and all sorts of metadata as well as the dependencies needed for building the app. As soon as all these were set, we started working on the three main categories of files mentioned above. Here is the simple structure we opted for.




Figure 6.1: Project Structure as shown in IntelliJ IDEA.

Architecturally speaking, because we have opted for a Component Based approach, everything is built so that it can be reused. Therefore, the nodes animators for all algorithms as well as the codeline animators are generic and only adapted to the specific algo they are running for.

Also, we have also added in assets, vendor and fonts components. These are just static files that will be used to render fonts, add images to the HTML pages as well as provide some extra JS modules to the specific algorithm animators.

Last but not least, the package.json file is automated so that we can add the releases as easy as possible to the release-builds folder which will contain all executables for the 3 main operating system families: macOS, Windows and Linux.

6.3 JavaScript and Multi-Threading

This was the most important issue faced throughout the whole software process. Because I did not have previous experience with web applications that might require multi-threading support, I did not know if JavaScript provides this or not. We shall provide a code snippet first to better describe the problem and discuss on it afterwards.

```
1 (function(ind) {
2     setTimeout(
3         function() {
4             unHighlightAllCodeLines();
5             highlightCodeLine(1);
6             if (nodes[ind] === nodeRoot) {
7                 distances[nodes[ind].label] = 0;
8             } else if (containsObject(nodes[ind], nodeRoot.adjacencyList)) {
9                 distances[nodes[ind].label] = getEdgeWeight(nodeRoot, nodes[ind]);
10            } else {
11                distances[nodes[ind].label] = Number.POSITIVE_INFINITY;
12            }
13            appendRowToTable(nodes[ind].label);
14            setupDistance(nodes[ind].label, distances[nodes[ind].label]);
15            if (ind > 0) {
16                if (!containsObject(nodes[ind - 1], S)) {
17                    nodes[ind - 1].color = "#009688";
18                }
19            }
20        }
21    );
22}
```

```

18     } else {
19         nodes[ind - 1].color = "#3f51b5";
20     }
21     unHighlightTableRow(nodes[ind - 1].label);
22 }
23 nodes[ind].color = "red";
24 highlightTableRow(nodes[ind].label);
25 network = rebuildNetwork(network, container, options, nodes);
26 },
27 2000 + 3000 * ind
28 );
29 })(i);

```

Listing 6.1: Animation function used in Dijkstra's SP algorithm.

As you can see above, this is a very complicated function which sets up initial distances to all nodes in a graph from the root - if they are adjacent, then the distance will be the weight of the edge; otherwise, the distance will be set to infinity. SetTimeout is a function that delays another function's execution until a specific point in time, thus we are sequencing multiple calls to this self-contained function based on the *i* index which is the variable used in the loop where this piece of code is defined.

This is the way all algorithm animations are implemented. However, the problem is as you can notice on lines 4, 5 or 21 that we are also animating other components apart from the nodes in the same body. In other programming languages, such as Java, this could easily be done by using 2 threads which would split the work between themselves - one would deal with node animation, while the other could highlight and unhighlight codelines dynamically. However, JavaScript does not allow support for multi-threading so we had to resume to merging the 2 functionalities into one. Even though one might argue that along with HTML5 the W3C introduced Web Workers [21], these do not provide a solution to our problem because they allow only primitives and simple objects to be passed inside the running thread, while we have Nodes that can have other Nodes as children (which is a property of the Vis Node), thus creating circular dependencies.

Therefore, this sort of complicated mathematical calculus had to be done (see line 27) in order to make sure the calls are invoked at the correct times and that they don't overlap.

Moreover, another type of issue occurred when there were setTimeout invocations inside outer timeouts. One particular problem we faced was that if you call a setTimeout function inside another setTimeout, the time you set the inner one to execute at is relative to the time the outer one will be invoked. Let's provide an example for a better visualization.

```

1 setTimeout(
2     function() {
3         unHighlightHeapCodeLine("delete", 1);
4         highlightHeapCodeLine("delete", 2);
5         impose(nodes[0]);
6     },
7     3000
8 );
9 // impose function starts here
10 for (let index = 0; index < steps; index++) {
11     (function(i) {
12         let prev = null;
13         highlightHeapCodeLine("impose", 1);
14         setTimeout(
15             function() {
16                 if (prev) {
17                     prev.color = "#009688";
18                 }
19                 prev = cursor;
20                 let largerValue;
21                 if (cursor.children.length === 1) {

```

```

22     largerValue = cursor.children[0];
23 } else {
24     if (cursor.children[0].label > cursor.children[1].label) {
25         largerValue = cursor.children[0];
26     } else {
27         largerValue = cursor.children[1];
28     }
29 }
30 cursor.color = "red";
31 largerValue.color = "red";
32 let temp = cursor.label;
33 cursor.label = largerValue.label;
34 largerValue.label = temp;
35 cursor = largerValue;
36 network = rebuildHeap(nodes, edges);
37 },
38 1000 * i
39 );
40 })(index);
41 }

```

Listing 6.2: Heap animation that uses nested setTimeouts.

The for loop below the comment is only a part of the impose function, but for shortening the length as much as possible we have included what was necessary to demonstrate the issue mentioned above.

As it can be observed, we have the call to the impose function (line 5) inside a setTimeout. Thus, that will execute after 3000 milliseconds, but inside the for loop every iteration will be invoked at 1000 milliseconds times index. Because of a lack of tutorials and resources on the web regarding such possible nested setTimeouts, I believed for a long time that the $1000 * i$ should be replaced with $3000 + 1000 * i$, but this was a big mistake because 3000 is inherited and it doesn't need to be included in the nested loop. After I realized that, the timings wouldn't overlap and the animation would perform normal.

6.4 Extra Features

Both the potential users and the supervisor have requested that all algorithms should contain a random functionality could easily have a graph or tree to animate that did not require manual insertion and deletion. Therefore, every algorithm page has a random button which will automate the whole process of creating a Vis network.

Also, for the Huffman encoding algorithm, we have also added an Input File button that can take a plain text file from the user's machine and, after parsing the document, execute the algorithm against that input.

Moreover, another feature implemented that was very popular was the landing page in the application - as the user opens the animator they will be presented with a carousel where they can interactively select what algorithm to see animated. Thus, we needed to place emphasis on this rather simple stylistic element so that users would enjoy the application since they opened it for the first time. Below a screenshot of the landing point is attached.



Figure 6.2: Landing page in the application.

Last but not least, we believe that having a modern and attractive webpage where users can download the app and get to know some basic information is essential, thus we have created a custom website where we host the releases and users can simply download and run Palgo with only a couple of clicks. This was a very nice feature voted for by many of our end users.

Figure 6.3: Palgo website hosted by GitHub.

6.5 Lessons Learned

One of the most important lessons learned was that web development, even though it might sound easy, is not like that at all. Because everything around us nowadays is web-based or it has something to do with it, there are a lot of technologies in various sub-fields that help developers solve certain problems. As discussed in the previous sections, we have struggled with many design decisions mainly because of the large palette of choices. Therefore, having to research and eventually learn many of them so that argumented comparisons could be made regarding which one is better was time consuming. However, once the developer possesses the basics, he/she can evolve and apply the fundamentals to learn any new technology much easier.

Another lesson learned was that one needs to take very good care to not fall into the trap of "premature optimization". As we started the project back in first semester, I was very keen on implementing a large number of UI features such as stylish buttons or attractive fade-ins when the window content changed. However, these early iterations were designed for setting up the basics and implementing the algorithms, as well as animate them. Learning that optimizing prematurely can lead to wasted effort, I have changed my approach and started prioritizing our user stories more effectively.

A third lesson I learned was that any JavaScript can fall in the "callback hell". This is such a common pattern among web developers that people created a website that specifically addresses this issue and gives advice on how to avoid it [18]. I have also fell in this trap but with advice taken from the website mentioned previously, I could easily refactor and follow the best practices promoted by industry leading developers.

Another vital aspect I've learned throughout our software process was the importance of code style in JavaScript. Even though other programming languages promote a robust and consistent code style as well (e.g. Java, C), JavaScript is the most popular technology used thus it has a very widespread reach. Therefore, many styling guides have been created so that new developers would find the code easier to read and, subsequently, to understand. The formatter I have used throughout the whole project is a tool called Prettier [25] and the guide I used for reference is the one Airbnb have open sourced on GitHub (it is now one of the most popular projects on the platform with almost 50k stars [1]). This will ensure that if the project will gain contributors in the future their accommodation period with the codebase will be smooth and fast.

Overall, I believe that this project was a real success and it was a great way to both learn new technologies as well as learn how to research into previous work and apply the best software engineering practices. I think that without following set rules, deadlines and good methodologies already proven successful in the industry, no project can ever succeed.

Chapter 7

Testing

7.1 Unit Testing

7.2 Integration Testing

7.3 Prototype Evaluation

7.4 Results

Chapter 8

Conclusions

8.1 Open Source

8.2 Project Roadmap

8.3 Final Thoughts

8.4 Acknowledgements

Bibliography

- [1] Airbnb. Airbnb JavaScript Code Style Guide. <https://github.com/airbnb/javascript>.
- [2] Base. Base Website. <http://getbase.org>.
- [3] Kent Beck. Embracing change with extreme programming. *Computer*, 32(10):70–77, 1999.
- [4] Grady Booch. *Object Oriented Design: With Applications*. Benjamin/Cummings, 1991.
- [5] Bharat Choudhary and Shantu K Rakesh. An Approach using Agile Method for Software Development. *Innovation and Challenges in Cyber Security (ICICCS-INBUSH)*, 2016.
- [6] CircleCI. CircleCI Website. <https://circleci.com>.
- [7] Dai Clegg and Richard Barker. *Case method fast-track: a RAD approach*. Addison-Wesley Longman Publishing Co., Inc, 1994.
- [8] Connextra. Connextra User Story 2001: ConnextraStoryCard. https://en.wikipedia.org/wiki/User_story#cite_note-connextra-4.
- [9] Sarah A. Douglas Cristopher D. Hundhausen and John T. Stasko. A Meta-Study of Algorithm Visualization Effectiveness. *Journal of Visual Languages and Computing*, 13:259–290, 2002.
- [10] Oxford Dictionaries. *Oxford English Dictionary*. OUP Oxford, 2012.
- [11] E. W. Dijkstra. A note on two problems in connexion with graphs. *Numerische Mathematik*, 1:269–271, 1959.
- [12] Electron. Electron Website. <https://electron.atom.io>.
- [13] David Galles. David Galles Algorithm Animator. <https://www.cs.usfca.edu/~galles/visualization/Algorithms.html>.
- [14] Google. Google Chrome V8 Engine Website. <https://developers.google.com/v8/>.
- [15] James Grenning. Planning poker or how to avoid analysis paralysis while release planning. *Renaissance Software Consulting*, 2002.
- [16] Simple Grid. Simple Grid Website. <http://thisisdallas.github.io/Simple-Grid>.
- [17] Steven et al. Halim. VisuAlgo Website. <http://visualgo.net>.
- [18] Callback Hell. Callback Hell Website. <http://callbackhell.com>.
- [19] Mike Mannion and Barry Keepence. SMART Requirements. *ACM SIGSOFT Software Engineering Notes*, 20(2):42–47, 1995.
- [20] John Morris. A toolkit for algorithm animation. *Proceedings AEESEAP*, 2004.
- [21] Mozilla. Mozilla Developer Network - Web Workers. https://developer.mozilla.org/en-US/docs/Web/API/Web_Workers_API.
- [22] NPM. NPM Website. <http://npmjs.com>.
- [23] Susan Palmiter and Jay Elkerton. An evaluation of animated demonstrations of learning computer-based tasks. *Proceedings of the SIGCHI conference on Human Factors in Computing Systems*, 1991.
- [24] Brenda Parker and Ian Mitchell. Effective methods for learning: a study in visualization. *Journal of Computing Sciences in Colleges*, 22(2):176–182, 2006.

- [25] Prettier. Prettier Github Page. <https://github.com/prettier/prettier>.
- [26] Markus Schüer Rößling, Guido and Bernd Freisleben. The ANIMAL algorithm animation tool. *ACM SIGCSE Bulletin*, 32(3), 2000.
- [27] Ian Sommerville. *Software Engineering*. Pearson, 2015.
- [28] StackOverflow. StackOverflow 2017 Survey. <http://stackoverflow.com/insights/survey/2017#technology>.
- [29] Matthias F Stallmann. Algorithm Animation with Galant. *IEEE Computer Graphics and Applications*, 37(1), 2017.
- [30] Trello. Trello Website. <https://trello.com>.
- [31] Velocity.js. Velocity.js Website. <http://velocityjs.org>.
- [32] VisJS. VisJs Website. <http://visjs.org>.
- [33] Federico Zivolo. Material Design Bootstrap Website. <http://fezvrasta.github.io/bootstrap-material-design>.